


```
import os
import zipfile
```

```
os.environ['KAGGLE_CONFIG_DIR'] = "/content"
!kaggle competitions download -c plant-pathology-2020-fgvc7
```


```
# Unzip the dataset
with zipfile.ZipFile("plant-pathology-2020-fgvc7.zip", 'r') as zip_ref:
    zip_ref.extractall("plant_data")
```

 Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /content/kaggle.json'

Downloading plant-pathology-2020-fgvc7.zip to /content  
 99% 774M/779M [00:04<00:00, 136MB/s]  
 100% 779M/779M [00:04<00:00, 201MB/s]

```
import pandas as pd
import numpy as np
import os
import cv2
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```


```
df = pd.read_csv("/content/plant_data/train.csv")
df.head()
```



	image_id	healthy	multiple_diseases	rust	scab
0	Train_0	0		0	1
1	Train_1	0		1	0
2	Train_2	1		0	0
3	Train_3	0		0	1
4	Train_4	1		0	0

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
df['label'] = df[['healthy', 'multiple_diseases', 'rust', 'scab']].idxmax(axis=1)
df = df[['image_id', 'label']]
df.head()
```



	image_id	label
0	Train_0	scab
1	Train_1	multiple_diseases
2	Train_2	healthy
3	Train_3	rust
4	Train_4	healthy

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
def extract_features(image_path):
    image = cv2.imread(image_path)
    image = cv2.resize(image, (128, 128)) # Resize for consistency
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

    hist_features = []
    for i in range(3): # RGB channels
        hist = cv2.calcHist([image], [i], None, [32], [0, 256])
        hist = cv2.normalize(hist, hist).flatten()
        hist_features.extend(hist)

    return hist_features

# Call the function with a valid image path
features = extract_features('/content/plant_data/images/Test_0.jpg')
```

```
print(features[:5]) # show first 5 features
```

```
↳ [np.float32(0.0), np.float32(0.002053484), np.float32(0.06853503), np.float32(0.3277874), np.float32(0.43995896)]
```

```
# Build feature Dataset
```

```
features = []
```

```
labels = []
```

```
image_folder = "/content/plant_data/images"
```

```
for i, row in df.iterrows():
```

```
    image_id = row["image_id"]
```

```
    label = row["label"]
```

```
    image_path = os.path.join(image_folder, image_id + '.jpg')
```

```
    feats = extract_features(image_path)
```

```
    features.append(feats)
```

```
    labels.append(label)
```

```
X = np.array(features)
```

```
y = np.array(labels)
```

```
# Encode Labels and Split data
```

```
le = LabelEncoder()
```

```
y_encoded = le.fit_transform(y)
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y_encoded, test_size=0.2, random_state=42, stratify=y_encoded
```

```
)
```

Double-click (or enter) to edit

```
# Train ML models
```

```
# A. Random Forest
```

```
rf = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
rf.fit(X_train, y_train)
```

```
rf_preds = rf.predict(X_test)
```

```
# B. SVM
```

```
svm = SVC(kernel = 'rbf', probability=True)
```

```
svm.fit(X_train, y_train)
```

```
svm_preds = svm.predict(X_test)
```

```
# C. Gradient Boosting
```

```
gbm = GradientBoostingClassifier(n_estimators=100)
```

```
gbm.fit(X_train, y_train)
```

```
gbm_preds = gbm.predict(X_test)
```

```
# Evaluate Models
```

```
def evaluate_model(y_true, y_pred, model_name):
```

```
    print(f"----- {model_name} Evaluation -----")
```

```
    print("Classification Report:\n", classification_report(y_true, y_pred, target_names=le.classes_))
```

```
    cm = confusion_matrix(y_true, y_pred)
```

```
    sns.heatmap(cm, annot=True, fmt='d', xticklabels=le.classes_, yticklabels=le.classes_, cmap='Blues')
```

```
    plt.title(f"{model_name} - Confusion Matrix")
```

```
    plt.xlabel('Predicted')
```

```
    plt.ylabel('True')
```

```
    plt.show()
```

```
evaluate_model(y_test, rf_preds, "Random Forest")
```

```
evaluate_model(y_test, svm_preds, "SVM")
```


```
evaluate_model(y_test, gbm_preds, "Gradient Boosting")
```

----- Random Forest Evaluation -----

Classification Report:

	precision	recall	f1-score	support
healthy	0.60	0.61	0.61	103
multiple_diseases	0.00	0.00	0.00	18
rust	0.64	0.66	0.65	125
scab	0.62	0.67	0.64	119
accuracy			0.62	365
macro avg	0.46	0.48	0.47	365
weighted avg	0.59	0.62	0.60	365

Random Forest - Confusion Matrix




----- SVM Evaluation -----

Classification Report:

	precision	recall	f1-score	support
healthy	0.57	0.52	0.55	103
multiple_diseases	0.00	0.00	0.00	18
rust	0.54	0.51	0.53	125
scab	0.51	0.65	0.57	119
accuracy			0.53	365
macro avg	0.40	0.42	0.41	365
weighted avg	0.51	0.53	0.52	365

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar  
\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar  
\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar  
\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))

SVM - Confusion Matrix




----- Gradient Boosting Evaluation -----

Classification Report:

	precision	recall	f1-score	support
healthy	0.57	0.58	0.58	103
multiple diseases	0.33	0.06	0.10	18

	rust	0.62	0.64	0.63	125
	scab	0.62	0.67	0.65	119
accuracy				0.61	365
macro avg	0.54	0.49	0.49		365
weighted avg	0.59	0.61	0.59		365



```
# 1. Train the Gradient Boosting model
from sklearn.ensemble import GradientBoostingClassifier

gbm = GradientBoostingClassifier()
gbm.fit(X_train, y_train)

# 2. Predict
y_pred_gbm = gbm.predict(X_test)

# 3. Visualize Classification Report as Bar Plot
from sklearn.metrics import classification_report
import pandas as pd
import matplotlib.pyplot as plt

# If you used LabelEncoder for y labels:
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
le.fit(y_train) # or y, if not split yet
labels = le.classes_


report = classification_report(y_test, y_pred_gbm, target_names=labels, output_dict=True)
df_report = pd.DataFrame(report).transpose()

# Filter class-wise metrics
df_classes = df_report.iloc[:-3][['precision', 'recall', 'f1-score']]

# Plot
df_classes.plot(kind='bar', figsize=(10, 6))
plt.title('Gradient Boosting - Class-wise Metrics')
plt.ylabel('Score')
plt.ylim(0, 1)
plt.grid(axis='y')
plt.xticks(rotation=45)
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()
```



Gradient Boosting - Class-wise Metrics




```
# Feature importnace plot
import numpy as np

# Assuming model is named `gbm`
importances = gbm.feature_importances_
indices = np.argsort(importances)[-20:] # top 20 important features

plt.figure(figsize=(10, 6))
plt.title('Top 20 Feature Importances - Gradient Boosting')
plt.barh(range(len(indices)), importances[indices], align='center')
plt.xlabel('Relative Importance')
plt.ylabel('Feature Index')
plt.grid()
plt.show()
```



Top 20 Feature Importances - Gradient Boosting



Start coding or [generate](#) with AI.